

Internship Report

A semantic-driven approach to system call filtering

ZEGAOUI Taquyedine

FiT building, Tsinghua University, Beijing, China

September 1, 2016



清华大学
Tsinghua University



POLYTECH
GRENOBLE

Table of contents

Context of project

Introducing JITK

- Overview of JITK

- System call policies enforcement: BPF

- User-friendly rules definition: SCPL

Introducing K-framework

- Overview of K-framework

- K-framework function and role

Design Choices

- Syntactic approaches

- Semantic approaches

Results and evolutions

Discussion

- Advantages

- Drawbacks

Section 1

Context of project

Internship Report

- └ Context of project

- └ Introducing JITK

beginframe

What do Bitcoin transaction scripting, network packet filtering and power management have in common ?

They all make use of in-kernel interpreters !

What do Bitcoin transaction scripting, network packet filtering and power management have in common ?
They all make use of in-kernel interpreters !

In-kernel interpreters

- ▶ In-kernel interpreters have become a staple of modern computation processes
- ▶ They also have become a major concern regarding security

In-kernel interpreters

- ▶ In-kernel interpreters have become a staple of modern computation processes
- ▶ They also have become a major concern regarding security
- ▶ Risks of malicious attacks or intern errors

In-kernel interpreters

- ▶ In-kernel interpreters have become a staple of modern computation processes
- ▶ They also have become a major concern regarding security
- ▶ Risks of malicious attacks or intern errors
- ▶ In-kernel interpreters run in kernel space

In-kernel interpreters

- ▶ In-kernel interpreters have become a staple of modern computation processes
- ▶ They also have become a major concern regarding security
- ▶ Risks of malicious attacks or intern errors
- ▶ In-kernel interpreters run in kernel space
- ▶ Any error or attack can have tremendous consequences

In-kernel interpreters

- ▶ In-kernel interpreters have become a staple of modern computation processes
- ▶ They also have become a major concern regarding security
- ▶ Risks of malicious attacks or intern errors
- ▶ In-kernel interpreters run in kernel space
- ▶ Any error or attack can have tremendous consequences

In-kernel interpreters

- ▶ In-kernel interpreters have become a staple of modern computation processes
- ▶ They also have become a major concern regarding security
- ▶ Risks of malicious attacks or intern errors
- ▶ In-kernel interpreters run in kernel space
- ▶ Any error or attack can have tremendous consequences

What can be done to guarantee safe and correct in-kernel interpreters ?

In-kernel interpreters

- ▶ In-kernel interpreters have become a staple of modern computation processes
- ▶ They also have become a major concern regarding security
- ▶ Risks of malicious attacks or intern errors
- ▶ In-kernel interpreters run in kernel space
- ▶ Any error or attack can have tremendous consequences

What can be done to guarantee safe and correct in-kernel interpreters ?

JITK: what is it ?

Enter JITK,

- ▶ An infrastructure for building in-kernel interpreters
- ▶ Guarantees complete functional correctness

JITK: what is it ?

Enter JITK,

- ▶ An infrastructure for building in-kernel interpreters
- ▶ Guarantees complete functional correctness
- ▶ Uses user-defined system call policies to secure execution

JITK: what is it ?

Enter JITK,

- ▶ An infrastructure for building in-kernel interpreters
- ▶ Guarantees complete functional correctness
- ▶ Uses user-defined system call policies to secure execution
- ▶ Those policies define what system call the executing code can invoke

JITK: what is it ?

Enter JITK,

- ▶ An infrastructure for building in-kernel interpreters
- ▶ Guarantees complete functional correctness
- ▶ Uses user-defined system call policies to secure execution
- ▶ Those policies define what system call the executing code can invoke

JITK: What does it do ?

The system call policies are

- ▶ Compiled to native code to be passed to the kernel
- ▶ Guaranteed to correctly translate into native code without loss of meaning

JITK: What does it do ?

The system call policies are

- ▶ Compiled to native code to be passed to the kernel
- ▶ Guaranteed to correctly translate into native code without loss of meaning
- ▶ Used to determine behavior regarding every system call

JITK: What does it do ?

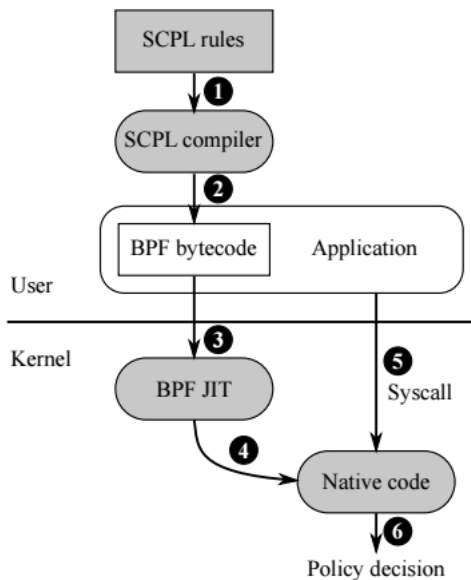
The system call policies are

- ▶ Compiled to native code to be passed to the kernel
- ▶ Guaranteed to correctly translate into native code without loss of meaning
- ▶ Used to determine behavior regarding every system call
- ▶ Able to be checked by the interpreter without overhead cost

JITK: What does it do ?

The system call policies are

- ▶ Compiled to native code to be passed to the kernel
- ▶ Guaranteed to correctly translate into native code without loss of meaning
- ▶ Used to determine behavior regarding every system call
- ▶ Able to be checked by the interpreter without overhead cost



Introducing BPF

The BPF language

- ▶ Originally serves for defining packets filters
- ▶ Is a low-level language rather close to assembly

Introducing BPF

The BPF language

- ▶ Originally serves for defining packets filters
- ▶ Is a low-level language rather close to assembly
- ▶ Is used here as a system call filter

Introducing BPF

The BPF language

- ▶ Originally serves for defining packets filters
- ▶ Is a low-level language rather close to assembly
- ▶ Is used here as a system call filter

Example of BPF

```
; load syscall number
ld [0]
; deny open() with errno = EACCES
jeq #SYS_open, L1, L2
L1: ret #RET_ERRNO|#EACCES
; allow getpid()
L2: jeq #SYS_getpid, L3, L4
L3: ret #RET_ALLOW
; allow gettimeofday()
L4: jeq #SYS_gettimeofday, L5, L6
L5: ret #RET_ALLOW
L6: ...
; default: kill current process
ret #RET_KILL
```

As seen above, each system call gets an entry in the list of rules, along with the expected behavior regarding this particular system call. A default behavior is also defined, should any system call be absent from the previous list.

Example of BPF

```
; load syscall number
ld [0]
; deny open() with errno = EACCES
jeq #SYS_open, L1, L2
L1: ret #RET_ERRNO|#EACCES
; allow getpid()
L2: jeq #SYS_getpid, L3, L4
L3: ret #RET_ALLOW
; allow gettimeofday()
L4: jeq #SYS_gettimeofday, L5, L6
L5: ret #RET_ALLOW
L6: ...
; default: kill current process
ret #RET_KILL
```

As seen above, each system call gets an entry in the list of rules, along with the expected behavior regarding this particular system call. A default behavior is also defined, should any system call be absent from the previous list.

Introducing SCPL

We introduce SCPL, a domain specific language for defining system call policies

- ▶ In a more user-friendly way, being close to natural human language
- ▶ In an easier, more intuitive and less prone to errors manner

Introducing SCPL

We introduce SCPL, a domain specific language for defining system call policies

- ▶ In a more user-friendly way, being close to natural human language
- ▶ In an easier, more intuitive and less prone to errors manner
- ▶ Which reduces the risk of having incorrect BPF policies

Introducing SCPL

We introduce SCPL, a domain specific language for defining system call policies

- ▶ In a more user-friendly way, being close to natural human language
- ▶ In an easier, more intuitive and less prone to errors manner
- ▶ Which reduces the risk of having incorrect BPF policies
- ▶ Which will be translated to BPF

Introducing SCPL

We introduce SCPL, a domain specific language for defining system call policies

- ▶ In a more user-friendly way, being close to natural human language
- ▶ In an easier, more intuitive and less prone to errors manner
- ▶ Which reduces the risk of having incorrect BPF policies
- ▶ Which will be translated to BPF

Example of SCPL

```
{ default_action = Kill;  
rules = [  
  { action = Errno EACCES; syscall = SYS_open };  
  { action = Allow; syscall = SYS_getpid };  
  { action = Allow; syscall = SYS_gettimeofday };  
  ...  
] }
```

As seen above, SCPL is really close to the natural thought process of defining the rules of system call behavior, and this intuitive ease of use guarantees minimal errors within policies definition.

Example of SCPL

```
{ default_action = Kill;  
rules = [  
  { action = Errno EACCES; syscall = SYS_open };  
  { action = Allow; syscall = SYS_getpid };  
  { action = Allow; syscall = SYS_gettimeofday };  
  ...  
] }
```

As seen above, SCPL is really close to the natural thought process of defining the rules of system call behavior, and this intuitive ease of use guarantees minimal errors within policies definition.

Languages implementations

- ▶ Too many languages have only vague semantics defined in incomplete manuals
- ▶ Several implementations can exist for a single language

Languages implementations

- ▶ Too many languages have only vague semantics defined in incomplete manuals
- ▶ Several implementations can exist for a single language
- ▶ Some or even all implementations often rely on ad-hoc intuition of said language semantics

Languages implementations

- ▶ Too many languages have only vague semantics defined in incomplete manuals
- ▶ Several implementations can exist for a single language
- ▶ Some or even all implementations often rely on ad-hoc intuition of said language semantics
- ▶ Results in formally wrong implementations, with possibly terrible results

Languages implementations

- ▶ Too many languages have only vague semantics defined in incomplete manuals
- ▶ Several implementations can exist for a single language
- ▶ Some or even all implementations often rely on ad-hoc intuition of said language semantics
- ▶ Results in formally wrong implementations, with possibly terrible results

Introducing K-framework

Enter K-framework, a semantic framework

- ▶ Allowing implementations to be strict applications of syntax and semantics
- ▶ Guaranteeing clean, ambiguity-free implementations conform to specifications

Introducing K-framework

Enter K-framework, a semantic framework

- ▶ Allowing implementations to be strict applications of syntax and semantics
- ▶ Guaranteeing clean, ambiguity-free implementations conform to specifications
- ▶ Eliminating errors stemming from approximative implementations

Introducing K-framework

Enter K-framework, a semantic framework

- ▶ Allowing implementations to be strict applications of syntax and semantics
- ▶ Guaranteeing clean, ambiguity-free implementations conform to specifications
- ▶ Eliminating errors stemming from approximative implementations

Producing formally proven implementations

K-framework is able to generate languages implementations

- ▶ Which correctness is guaranteed by the use of Coq and OCaml
- ▶ Completely trustworthy and free of errors

Producing formally proven implementations

K-framework is able to generate languages implementations

- ▶ Which correctness is guaranteed by the use of Coq and OCaml
- ▶ Completely trustworthy and free of errors
- ▶ Which can be used to define compilers between defined languages

Producing formally proven implementations

K-framework is able to generate languages implementations

- ▶ Which correctness is guaranteed by the use of Coq and OCaml
- ▶ Completely trustworthy and free of errors
- ▶ Which can be used to define compilers between defined languages

Only by defining syntax and semantics

To generate such language implementations, we only need to define its syntax and semantics, split into three components

- ▶ Configurations: descriptions of the program state using nested cells
- ▶ Computations: sequences of computational tasks as nested lists

Only by defining syntax and semantics

To generate such language implementations, we only need to define its syntax and semantics, split into three components

- ▶ Configurations: descriptions of the program state using nested cells
- ▶ Computations: sequences of computational tasks as nested lists
- ▶ Rules are a generalisation of conventional rewrite rules

Only by defining syntax and semantics

To generate such language implementations, we only need to define its syntax and semantics, split into three components

- ▶ Configurations: descriptions of the program state using nested cells
- ▶ Computations: sequences of computational tasks as nested lists
- ▶ Rules are a generalisation of conventional rewrite rules

Section 2

Design Choices

SCPL syntactic approach

Regarding the SCPL syntax

- ▶ DSL made for JITK
- ▶ Simple and close to thought process

SCPL syntactic approach

Regarding the SCPL syntax

- ▶ DSL made for JITK
- ▶ Simple and close to thought process
- ▶ No formal language definition

SCPL syntactic approach

Regarding the SCPL syntax

- ▶ DSL made for JITK
- ▶ Simple and close to thought process
- ▶ No formal language definition
- ▶ Ad-hoc definition used in JITK

SCPL syntactic approach

Regarding the SCPL syntax

- ▶ DSL made for JITK
- ▶ Simple and close to thought process
- ▶ No formal language definition
- ▶ Ad-hoc definition used in JITK
- ▶ Clean redefinition of the syntax

SCPL syntactic approach

Regarding the SCPL syntax

- ▶ DSL made for JITK
- ▶ Simple and close to thought process
- ▶ No formal language definition
- ▶ Ad-hoc definition used in JITK
- ▶ Clean redefinition of the syntax

BPF syntactic approach

Regarding the BPF syntax

- ▶ Widely used general purpose language
- ▶ Specific use as call filter

BPF syntactic approach

Regarding the BPF syntax

- ▶ Widely used general purpose language
- ▶ Specific use as call filter
- ▶ Numerous features left unused

BPF syntactic approach

Regarding the BPF syntax

- ▶ Widely used general purpose language
- ▶ Specific use as call filter
- ▶ Numerous features left unused
- ▶ Redefinition of relevant used syntax

BPF syntactic approach

Regarding the BPF syntax

- ▶ Widely used general purpose language
- ▶ Specific use as call filter
- ▶ Numerous features left unused
- ▶ Redefinition of relevant used syntax

Comparative waterfall approach

Rewriting-based semantic definition

- ▶ User input is successively compared to each rule
- ▶ A rule with a different system call gets deleted

Comparative waterfall approach

Rewriting-based semantic definition

- ▶ User input is successively compared to each rule
- ▶ A rule with a different system call gets deleted
- ▶ A matching rule deletes all other rules but itself

Comparative waterfall approach

Rewriting-based semantic definition

- ▶ User input is successively compared to each rule
- ▶ A rule with a different system call gets deleted
- ▶ A matching rule deletes all other rules but itself
- ▶ Result is extracted from the last standing rule

Comparative waterfall approach

Rewriting-based semantic definition

- ▶ User input is successively compared to each rule
- ▶ A rule with a different system call gets deleted
- ▶ A matching rule deletes all other rules but itself
- ▶ Result is extracted from the last standing rule

Variable map approach

Environment-based semantic definition

- ▶ Each rule is first translated into a map variable
- ▶ Variable couple: `SystemCall` bound to `ActionCode`

Variable map approach

Environment-based semantic definition

- ▶ Each rule is first translated into a map variable
- ▶ Variable couple: `SystemCall` bound to `ActionCode`
- ▶ Comparison is treated as a variable lookup

Variable map approach

Environment-based semantic definition

- ▶ Each rule is first translated into a map variable
- ▶ Variable couple: `SystemCall` bound to `ActionCode`
- ▶ Comparison is treated as a variable lookup
- ▶ Result is the value of the matching name variable

Variable map approach

Environment-based semantic definition

- ▶ Each rule is first translated into a map variable
- ▶ Variable couple: `SystemCall` bound to `ActionCode`
- ▶ Comparison is treated as a variable lookup
- ▶ Result is the value of the matching name variable

Abstract assembly approach

Use of an abstract machine for BPF semantic

- ▶ Being an assembly language, BPF relies on machine structures
- ▶ Abstraction of those structures through the use of K own structures

Abstract assembly approach

Use of an abstract machine for BPF semantic

- ▶ Being an assembly language, BPF relies on machine structures
- ▶ Abstraction of those structures through the use of K own structures
- ▶ Abstract configuration cells tasked with machine simulation

Abstract assembly approach

Use of an abstract machine for BPF semantic

- ▶ Being an assembly language, BPF relies on machine structures
- ▶ Abstraction of those structures through the use of K own structures
- ▶ Abstract configuration cells tasked with machine simulation

Section 3

Results and evolutions

Final features

Final features

- ▶ Full SCPL and BPF syntaxes, true to the new definitions
- ▶ Parsing of user-submitted SCPL or BPF scripts

Final features

Final features

- ▶ Full SCPL and BPF syntaxes, true to the new definitions
- ▶ Parsing of user-submitted SCPL or BPF scripts
- ▶ Evaluation and splitting of scripts in computation cells

Final features

Final features

- ▶ Full SCPL and BPF syntaxes, true to the new definitions
- ▶ Parsing of user-submitted SCPL or BPF scripts
- ▶ Evaluation and splitting of scripts in computation cells
- ▶ Waterfall and environment based comparisons in BPF

Final features

Final features

- ▶ Full SCPL and BPF syntaxes, true to the new definitions
- ▶ Parsing of user-submitted SCPL or BPF scripts
- ▶ Evaluation and splitting of scripts in computation cells
- ▶ Waterfall and environment based comparisons in BPF

Delayed features

Delayed features

- ▶ Full machine abstraction via configuration cells
- ▶ Proper syntactic type casting

Delayed features

Delayed features

- ▶ Full machine abstraction via configuration cells
- ▶ Proper syntactic type casting
- ▶ Proper input and output behavior

Delayed features

Delayed features

- ▶ Full machine abstraction via configuration cells
- ▶ Proper syntactic type casting
- ▶ Proper input and output behavior

Possible evolutions

Possible evolutions

- ▶ Formal SCPL and BPF syntax and semantic definitions
- ▶ K-framework guide to make up for rare and rough online documentation

Possible evolutions

Possible evolutions

- ▶ Formal SCPL and BPF syntax and semantic definitions
- ▶ K-framework guide to make up for rare and rough online documentation
- ▶ Proper logical proof research

Possible evolutions

Possible evolutions

- ▶ Formal SCPL and BPF syntax and semantic definitions
- ▶ K-framework guide to make up for rare and rough online documentation
- ▶ Proper logical proof research
- ▶ SCPL to BPF compiler

Possible evolutions

Possible evolutions

- ▶ Formal SCPL and BPF syntax and semantic definitions
- ▶ K-framework guide to make up for rare and rough online documentation
- ▶ Proper logical proof research
- ▶ SCPL to BPF compiler

Section 4

Discussion

Advantages of K

- ▶ Lightweight, portable implementations
- ▶ Relative safety and correctness provided

Advantages of K

- ▶ Lightweight, portable implementations
- ▶ Relative safety and correctness provided
- ▶ Deterministic and explicit behavior

Advantages of K

- ▶ Lightweight, portable implementations
- ▶ Relative safety and correctness provided
- ▶ Deterministic and explicit behavior
- ▶ Relative ease and efficiency of implementation once used

Advantages of K

- ▶ Lightweight, portable implementations
- ▶ Relative safety and correctness provided
- ▶ Deterministic and explicit behavior
- ▶ Relative ease and efficiency of implementation once used

Drawbacks of K

Coq

- ▶ Rapidly changing project, esp parsers
- ▶ Learning curve for implementing languages

Drawbacks of K

Coq

- ▶ Rapidly changing project, esp parsers
- ▶ Learning curve for implementing languages
- ▶ Lack of definite and up-to-date documentation and ressources

Drawbacks of K

Coq

- ▶ Rapidly changing project, esp parsers
- ▶ Learning curve for implementing languages
- ▶ Lack of definite and up-to-date documentation and ressources
- ▶ Somewhat lacking safety and correctness compared to Coq

Drawbacks of K

Coq

- ▶ Rapidly changing project, esp parsers
- ▶ Learning curve for implementing languages
- ▶ Lack of definite and up-to-date documentation and ressources
- ▶ Somewhat lacking safety and correctness compared to Coq

Section 5

Thank you for your attention ! Any questions ?

Further Reading I



Xi Wang, David Lazar, Nickolai Zeldovich, Adam Chlipala,
Zachary Tatlock

Jitk: A Trustworthy In-Kernel Interpreter Infrastructure.
11th USENIX Symposium on Operating Systems Design and
Implementation, 2014.



Grigore Rosu, Traian Florin Serbanuta

An Overview of the K Semantic Framework
J.LAP, 2010.



Steven McCanne and Van Jacobson Lawrence Berkeley
Laboratory

The BSD Packet Filter: A New Architecture for User-level
Packet Capture
December 19, 1992

Further Reading II



Mihail Asvoae, Computer Science Department, University "Al I
Cuza" Iasi, Romania

K Semantics for Assembly Languages: A Case Study

Electronic Notes in Theoretical Computer Science 304 (2014)
111125